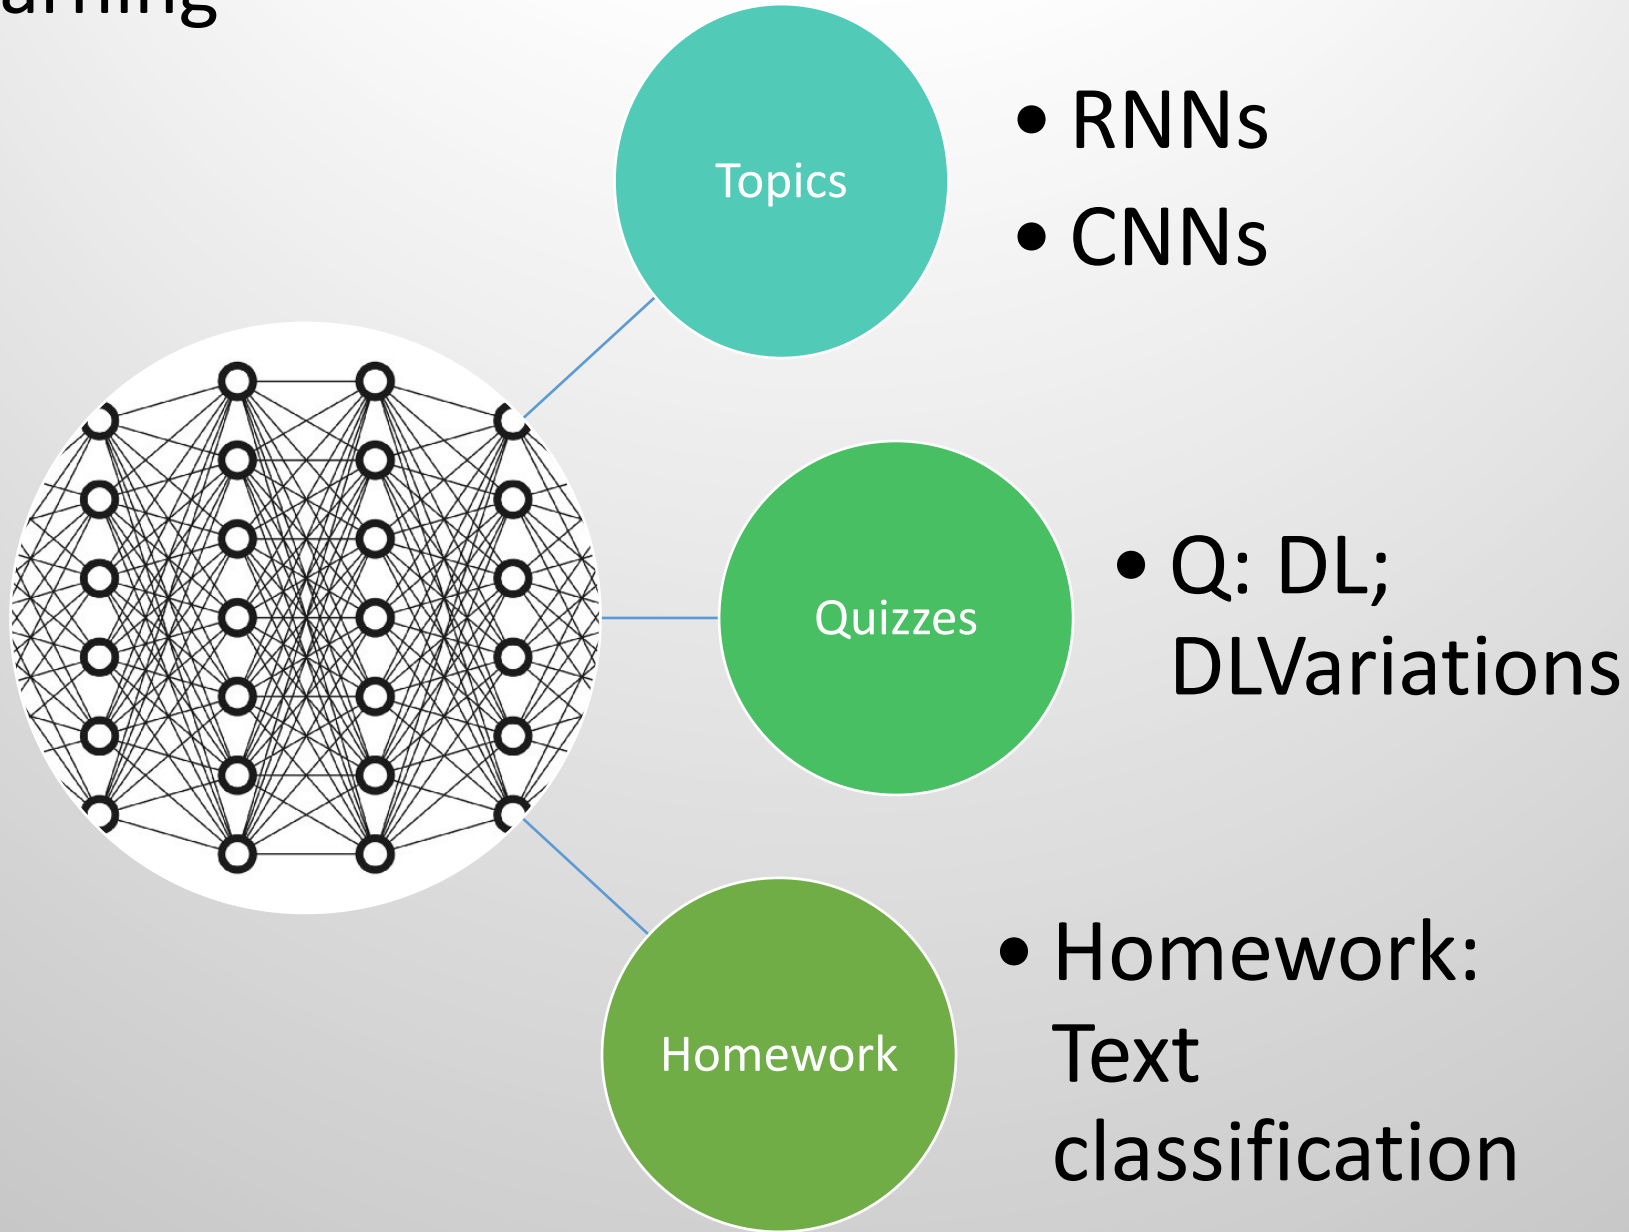


Natural Language Processing

Dr. Karen Mazidi



Part Six: Deep Learning



CNNs

- convolutional neural networks or covnets
- densely connected sequential layers learn global patterns in the data
- CNNs learn patterns in small windows
- Advantages:
 - a pattern learned in one location is recognized elsewhere
 - layers can learn hierarchies of shapes from edges and other features

CNN convolution

- a 4x4 'filter' slides over the data, performing the convolution function

0.47557	0.13031	0.26269	0.98775	0.54559	0.70388	0.41101	0.10889	
0.25782	0.69232	0.53866	0.20306	0.01652	0.45732	0.49489	0.47130	
0.87015	0.03241	0.00089	0.95473	0.25201	0.67926	0.66318	0.35740	
0.13696	0.20884	0.20363	0.72029	0.26433	0.42732	0.87660	0.59141	
0.51279	0.81518	0.50046	0.89543	0.77181	0.77192	0.45861	0.25983	
0.03777	0.12560	0.54588	0.06574	0.31243	0.50573	0.60777	0.85029	
0.82038	0.42600	0.16205	0.80647	0.10582	0.45355	0.59760	0.08356	
0.71715	0.42875	0.85921	0.60168	0.92237	0.62636	0.71523	0.14542	
0.09399	0.43249	0.84148	0.23740	0.30299	0.93350	0.03851	0.33104	
0.30386	0.63560	0.72024	0.38294	0.78565	0.72367	0.52017	0.93030	
0.97332	0.02479	0.31189	0.74439	0.62472	0.62113	0.13827	0.92139	
0.95440	0.03045	0.41130	0.71335	0.07405	0.03085	0.43504	0.93417	

Figure 24.1: Convolving

convolution

- convolution – mathematical process of combining two functions
- the filter (aka kernel) moves with overlap in strides, the smaller the stride, the more the overlap

2	4	9	1	4
2	1	4	4	6
1	1	2	9	2
7	3	5	1	3
2	3	4	8	5

Image

x

1	2	3
-4	7	4
2	-5	1

Filter /
Kernel

=

51		

Feature

padding

- notice the shrinkage of the data, at least one per dimension
- this can be avoided by padding the data

2	4	9	1	4
2	1	4	4	6
1	1	2	9	2
7	3	5	1	3
2	3	4	8	5

Image

x

1	2	3
-4	7	4
2	-5	1

Filter /
Kernel

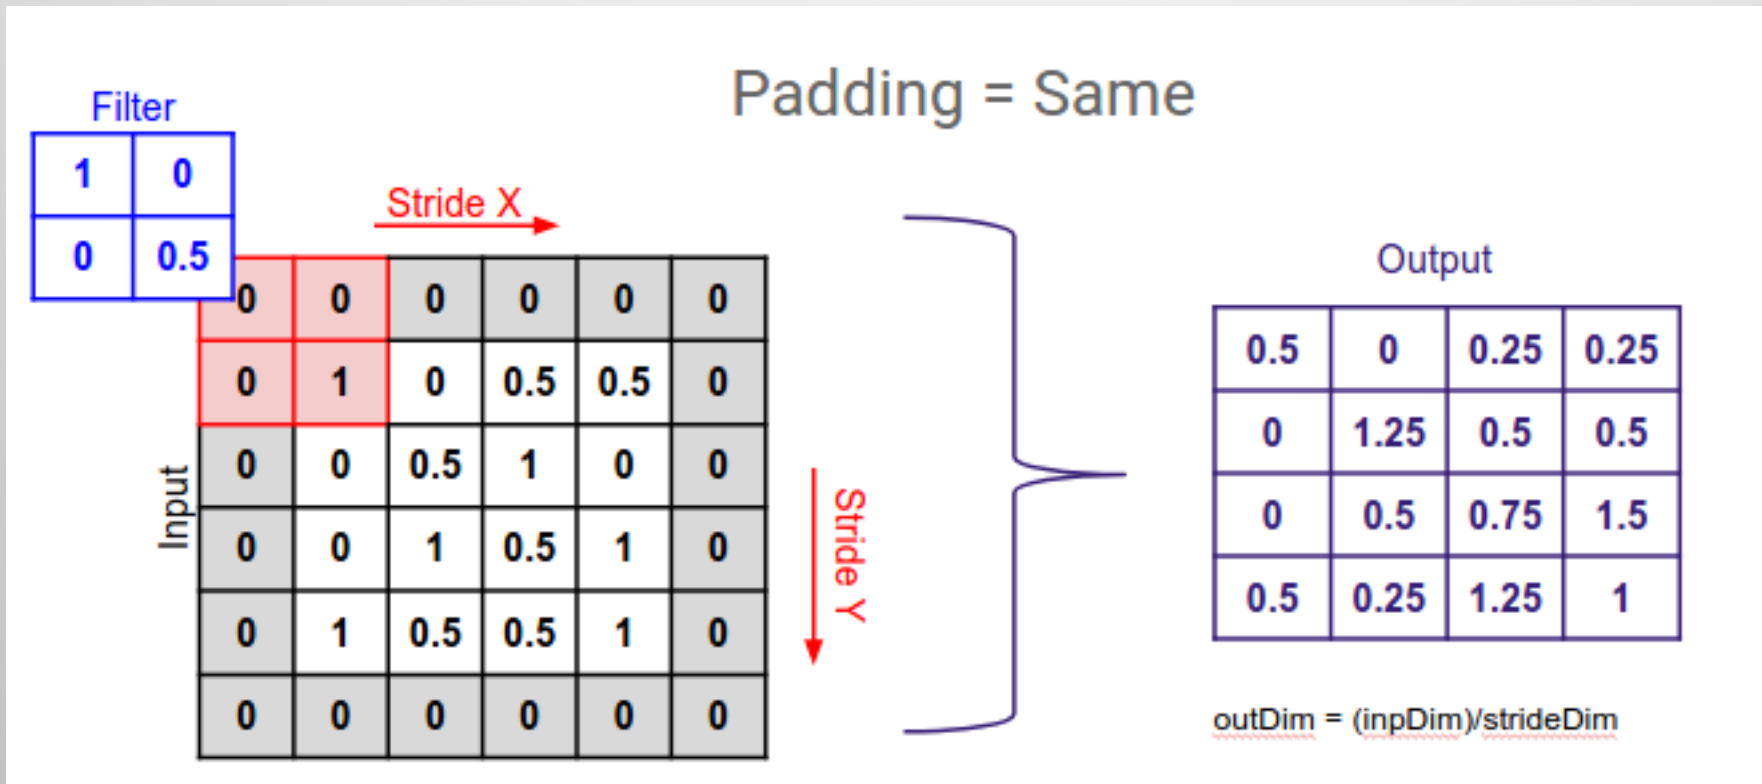
=

51		

Feature

padding

- padding = same gives same output size of data
- padding = none doesn't do padding




CNN

- Conv1D layers work well on text data
- stacks of conv layers and max-pooling layers are common, followed by a flatten layer, then a dense layer for the final classification
- max pooling also reduces dimensions and helps prevent overfitting

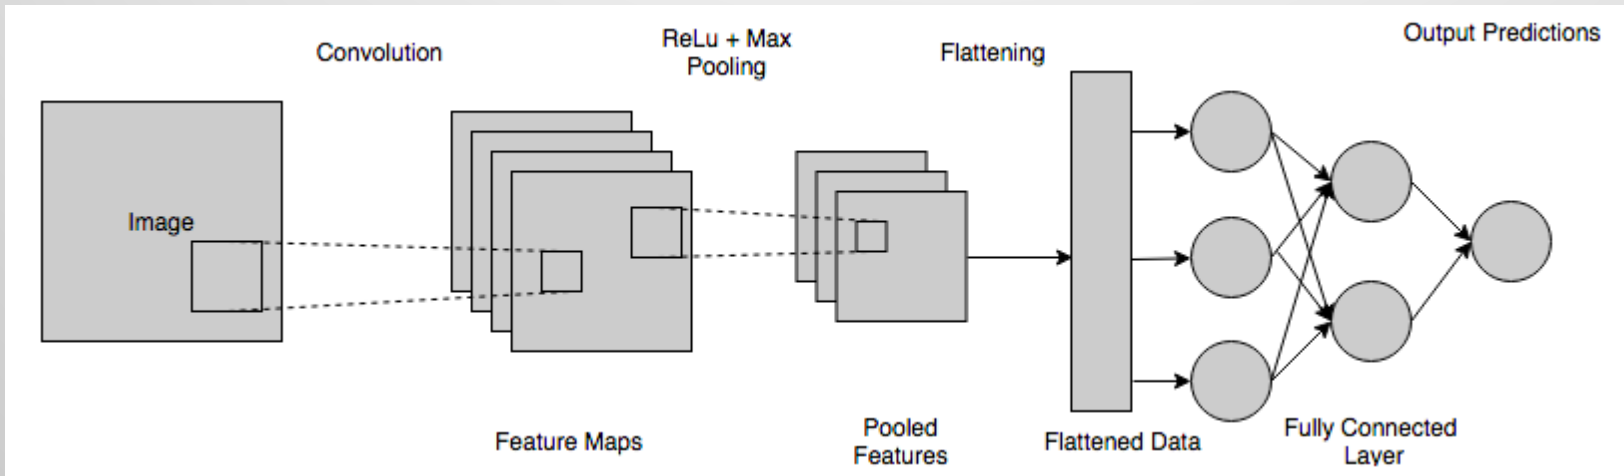
max pooling

- dimensionality reduction



Flatten

- after conv-max pooling layers, flattening reshapes the data for further processing



CNN visualization

- https://www.youtube.com/watch?v=YRhxdVk_sIs

Keras: IMDB data

- each example was shortened or padded to length 500
- input shape is (25000, 500)
- the embedding layer learns connections between words

```
model = models.Sequential()  
model.add(layers.Embedding(max_features, 128, input_length=maxlen))  
model.add(layers.Conv1D(32, 7, activation='relu'))  
model.add(layers.MaxPooling1D(5))  
model.add(layers.Conv1D(32, 7, activation='relu'))  
model.add(layers.GlobalMaxPooling1D())  
model.add(layers.Dense(1))
```

Keras: IMDB data

Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 500, 128)	1280000
conv1d_2 (Conv1D)	(None, 494, 32)	28704
max_pooling1d_1 (MaxPooling1D)	(None, 98, 32)	0
conv1d_3 (Conv1D)	(None, 92, 32)	7200
global_max_pooling1d_1 (GlobalMaxPooling1D)	(None, 32)	0
dense_1 (Dense)	(None, 1)	33

Total params: 1,315,937

Trainable params: 1,315,937

Non-trainable params: 0


Keras: IMDB data

- train and test as before
- results: a couple of points higher than the sequential model

Recurrent models

- a recurrent neural network, RNN, has memory, or state, which enables it to learn a sequence
- the looping mechanism produces a new hidden state at each iteration
- the final hidden state is a representation of previous states

RNN showing loop



Unrolled RNN




Figure 24.3: Recurrent Neural Network

RNNs

- vanishing gradient problem: with more layers, the back-propagated gradient becomes smaller and smaller
- LSTM (Long Short-Term Memory) is an RNN variation that helps the vanishing gradient problem
- keeps memory path independent of the back prop path
- LSTM allows information to be remembered or forgotten
- GRU is simpler than LSTM and may train faster

LSTM and GRU visualization

- <https://www.youtube.com/watch?v=8HyCNIVRbSU>

RNN on IMDB data

```
model = models.Sequential()
model.add(layers.Embedding(max_features, 32))
model.add(layers.SimpleRNN(32))
model.add(layers.Dense(1, activation='sigmoid'))
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
=====		
embedding_3 (Embedding)	(None, None, 32)	320000
=====		
simple_rnn_1 (SimpleRNN)	(None, 32)	2080
=====		
dense_2 (Dense)	(None, 1)	33
=====		

Total params: 322,113

Trainable params: 322,113

Non-trainable params: 0

LSTM on IMDB data

```
# build a model with LSTM
model = models.Sequential()
model.add(layers.Embedding(max_features, 32))
model.add(layers.LSTM(32))
model.add(layers.Dense(1, activation='sigmoid'))
```

GRU on IMDB data

```
model = models.Sequential()  
model.add(layers.Embedding(max_features, 32))  
model.add(layers.GRU(32))  
model.add(layers.Dense(1, activation='sigmoid'))
```


Code Examples

- Keras imdb 2 with RNN
- Keras imdb 3 with CNN



Essential points to note

- CNNs perform well on image data but also work on text data
- RNNs were created for sequential data like text, but suffer from vanishing gradients
- LSTM and GRU are improvements over the RNN

To Do

- Quiz on deep learning variations
- Portfolio: Text classification

The image shows a 'TO DO' list template with a blue header and a white body. The header contains the title 'TO DO' in large white letters, and three input fields for 'DATE:', 'FINISH BY:', and 'TOPIC:'. The body is divided into four sections: 'TASKS', 'ERRANDS', 'CORRESPONDENCE', and 'NOTES'. Each section has a 'No.' column and a 'DONE' column. The 'TASKS' and 'ERRANDS' sections have 10 rows each, and the 'CORRESPONDENCE' and 'NOTES' sections have 10 rows each. At the bottom right, there is a checkbox labeled 'ALL DONE'. At the bottom left, there is a quote: 'Make a list—you'll feel better.'.

TASKS		DONE	ERRANDS		DONE
No.					
01					
02					
03					
04					
05					
06					
07					
08					
09					
10					

CORRESPONDENCE		DONE	NOTES		DONE
No.					
01					
02					
03					
04					
05					
06					
07					
08					
09					
10					

ALL DONE

"Make a list—you'll feel better."

Next topic

embeddings

